

Experiment No:01

Aim:Data preparation using NumPy and Pandas

Theory:

Data Preprocessing:

Definition:

Data preprocessing is a crucial step in the data analysis and machine learning pipeline. It involves cleaning and transforming raw data into a format suitable for analysis or training machine learning models. The process includes handling missing values, removing duplicates, scaling, encoding categorical variables, and addressing outliers.

Purpose:

Data preprocessing enhances the quality and reliability of the data, making it suitable for analysis and model training. It helps to mitigate the impact of noise and inconsistencies in the data, leading to more accurate and robust results.

Numpy:

NumPy (Numerical Python) is a powerful library for numerical computing in Python. It provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on these arrays efficiently.

Key Features:

- Efficient array operations and mathematical functions.
- Broadcasting to perform operations on arrays of different shapes.
- Integration with other data analysis libraries like Pandas.

Use Cases:

NumPy is extensively used in scientific computing, machine learning, and data analysis for its array handling capabilities and numerical operations.

Pandas:

Pandas is a data manipulation and analysis library for Python. It provides data structures like DataFrame and Series, which are powerful tools for handling and analyzing structured data.

Key Features:

- DataFrame and Series for easy data manipulation.
- Built-in functions for data cleaning, exploration, and transformation.
- Integration with other libraries like NumPy and scikit-learn.

Use Cases:

Pandas is widely used for data cleaning, exploration, and preprocessing tasks in data science and analysis projects.

Normalization:

Normalization is the process of scaling numerical features to a standard range, often between 0 and 1. This ensures that all features contribute equally to the analysis or model training. In scikit-learn, the StandardScaler is commonly used for normalization.

Commands:

1.Importing Libraries and Packages :

```
[1] import numpy as np
import pandas as pd
```

2.Load dataset into Pandas:

pd.read_csv()

The following command will load the data in pandas and will show us some rows and columns from our dataset.

```
df=pd.read_csv("jobsdata.csv")
```

3.Description of the dataset.

df.info()

This method prints information about a DataFrame including the index data type and columns, non-null values and memory usage.

```
[23] df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9355 entries, 0 to 9354
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   work_year              9355 non-null   int64  
 1   job_title              9355 non-null   object  
 2   job_category           9355 non-null   object  
 3   salary_currency        9355 non-null   object  
 4   salary                 9355 non-null   int64  
 5   salary_in_usd          9355 non-null   int64  
 6   employee_residence     9355 non-null   object  
 7   experience_level       9355 non-null   object  
 8   employment_type        9355 non-null   object  
 9   work_setting           9355 non-null   object  
10   company_location       9355 non-null   object  
11   company_size           9355 non-null   object  
dtypes: int64(3), object(9)
memory usage: 877.2+ KB
```

4.Drop column that are not useful:

df.drop():

The drop() function is used to drop specified labels from rows or columns. Remove rows or columns by specifying label names and corresponding axis, or by specifying directly index or column names.

labels >> Index or column labels to drop.

axis >> Whether to drop labels from the index (0 or 'index') or columns (1 or 'columns').

```
[25] cols = ['employee_residence','company_size','work_setting']  
df =df.drop(cols,axis=1)
```

```
[26] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 9355 entries, 0 to 9354  
Data columns (total 9 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   work_year              9355 non-null   int64  
1   job_title              9355 non-null   object  
2   job_category           9355 non-null   object  
3   salary_currency        9355 non-null   object  
4   salary                 9355 non-null   int64  
5   salary_in_usd          9355 non-null   int64  
6   experience_level        9355 non-null   object  
7   employment_type        9355 non-null   object  
8   company_location       9355 non-null   object  
dtypes: int64(3), object(6)  
memory usage: 657.9+ KB
```

5.Take Care of missing values:

Let's compute a median or interpolate() the 'Código ISO del país' and fill those missing values. Pandas has an interpolate() function that will replace all the missing NaNs to interpolated values.

```
[25] cols = ['employee_residence','company_size','work_setting']  
df =df.drop(cols,axis=1)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 9355 entries, 0 to 9354  
Data columns (total 9 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   work_year              9355 non-null   int64  
1   job_title              9355 non-null   object  
2   job_category           9355 non-null   object  
3   salary_currency        9355 non-null   object  
4   salary                 9355 non-null   int64  
5   salary_in_usd          9355 non-null   int64  
6   experience_level        9355 non-null   object  
7   employment_type        9355 non-null   object  
8   company_location       9355 non-null   object  
dtypes: int64(3), object(6)  
memory usage: 657.9+ KB
```

6.Create Dummy Variables:

`pandas.get_dummies()` is used for data manipulation. It converts categorical data into dummy or indicator variables.

```
[30] dummies = []
      cols = ['experience_level', 'employment_type']
      for col in cols:
          dummies.append(pd.get_dummies(df[col]))

[31] temp =pd.concat(dummies,axis=1)

df = pd.concat((df,temp),axis=1)

[33] df = df.drop(['experience_level', 'employment_type'],axis=1)

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9355 entries, 0 to 9354
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   work_year              9355 non-null   int64
1   job_title              9355 non-null   object
2   job_category           9355 non-null   object
3   salary_currency        9355 non-null   object
4   salary                 9355 non-null   int64
5   salary_in_usd          9355 non-null   int64
6   company_location       9355 non-null   object
7   Entry-level            9355 non-null   uint8
8   Executive              9355 non-null   uint8
9   Mid-level              9355 non-null   uint8
10  Senior                 9355 non-null   uint8
11  Contract               9355 non-null   uint8
12  Freelance              9355 non-null   uint8
13  Full-time              9355 non-null   uint8
14  Part-time              9355 non-null   uint8
dtypes: int64(3), object(4), uint8(8)
memory usage: 584.8+ KB
```

7.Finding Outliers Manually:

In simple terms, an outlier is an extremely high or extremely low data point relative to the nearest data point and the rest of the neighboring co-existing values in a data graph or dataset you're working with. Outliers can give helpful insights into the data you're studying, and they can have an effect on statistical results. This can potentially help you discover inconsistencies and detect any errors in your statistical processes.

8.Normalization :

In statistics and machine learning, min-max normalization of data is a process of converting the original range of data to the range between 0 and 1. The resulting normalized values represent the original data on 0-1 scale. This will allow us to compare multiple features together and get more relevant information since now all the data will be on the same scale.

i) Using min-max :

```
[39] from sklearn.preprocessing import MinMaxScaler

[40] scaler = MinMaxScaler()
      for column in numerical_columns:
          df[column]= scaler.fit_transform(df[[column]])

[41] print(df.head())
```

	work_year		job_title		job_category	\
0	2023	Data	DevOps Engineer		Data Engineering	
1	2023		Data Architect	Data	Architecture and Modeling	
2	2023		Data Architect	Data	Architecture and Modeling	
3	2023		Data Scientist		Data Science and Research	
4	2023		Data Scientist		Data Science and Research	

	salary_currency	salary	salary_in_usd	company_location	Entry-level	\
0	EUR	88000	0.183936	Germany	0	
1	USD	186000	0.393103	United States	0	
2	USD	81800	0.153563	United States	0	
3	USD	212000	0.452874	United States	0	
4	USD	93300	0.180000	United States	0	

	Executive	Mid-level	Senior	Contract	Freelance	Full-time	Part-time	
0	0	1	0	0	0	1	0	
1	0	0	1	0	0	1	0	
2	0	0	1	0	0	1	0	
3	0	0	1	0	0	1	0	
4	0	0	1	0	0	1	0	

ii) Using Mean:

```
[36] numerical_columns=['salary_in_usd']

[37] for column in numerical_columns:
      mean_value = df[column].mean()
      df[column] = (df[column]-mean_value) / df[column].std()

      print(df.head())
```

	work_year		job_title		job_category	\
0	2023	Data	DevOps Engineer		Data Engineering	
1	2023		Data Architect	Data	Architecture and Modeling	
2	2023		Data Architect	Data	Architecture and Modeling	
3	2023		Data Scientist		Data Science and Research	
4	2023		Data Scientist		Data Science and Research	

	salary_currency	salary	salary_in_usd	company_location	Entry-level	\
0	EUR	88000	-0.875115	Germany	0	
1	USD	186000	0.565084	United States	0	
2	USD	81800	-1.084241	United States	0	
3	USD	212000	0.976623	United States	0	
4	USD	93300	-0.902214	United States	0	

	Executive	Mid-level	Senior	Contract	Freelance	Full-time	Part-time	
0	0	1	0	0	0	1	0	
1	0	0	1	0	0	1	0	
2	0	0	1	0	0	1	0	
3	0	0	1	0	0	1	0	
4	0	0	1	0	0	1	0	

Standardization:

In statistics and machine learning, data standardization is a process of converting data to z-score values based on the mean and standard deviation of the data. The resulting standardized value shows the number of standard deviations the raw value is away from the mean. Basically each value of a given feature of a dataset will be converted to a representative number of standard deviations that it's away from the mean of the feature.

```
from sklearn.preprocessing import StandardScaler
import pandas as pd

numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns

scaler = StandardScaler()
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])

print(df)
```

	work_year	job_title	job_category	\		
0	0.461170	Data DevOps Engineer	Data Engineering			
1	0.461170	Data Architect	Data Architecture and Modeling			
2	0.461170	Data Architect	Data Architecture and Modeling			
3	0.461170	Data Scientist	Data Science and Research			
4	0.461170	Data Scientist	Data Science and Research			
...			
9350	-3.389115	Data Specialist	Data Management and Strategy			
9351	-5.314258	Data Scientist	Data Science and Research			
9352	-3.389115	Principal Data Scientist	Data Science and Research			
9353	-5.314258	Data Scientist	Data Science and Research			
9354	-5.314258	Business Data Analyst	Data Analysis			
	salary_currency	salary	salary_in_usd	company_location	Entry-level	\
0	EUR	-0.973627	-0.875162	Germany	0	
1	USD	0.567122	0.565114	United States	0	
2	USD	-1.071103	-1.084299	United States	0	
3	USD	0.975892	0.976676	United States	0	
4	USD	-0.890301	-0.902262	United States	0	
...	

Conclusion: Thus we have implemented the concepts of data normalization on the dataset of job_data.